

Predicting Dielectric Properties for 49,566 Materials at 99% Confidence

Bayesian Cross-Categorization Replaces Expensive DFPT Calculations
with Calibrated Predictions from Structural Features Alone

JAX-CrossCat (jaxcross)

<https://github.com/sambhal-labs/jaxcross>

Contributions

- 49,566 materials with predicted ionic dielectric at 99% CI (publishable dataset)
- $R^2=0.81$ on holdout validation | 99% CI achieves 99.7% coverage
- 5 physically meaningful property groups discovered (phonon vs band structure)
- 88% of Random Forest accuracy from a single unsupervised model
- DFPT calculations cost 5-10x standard DFT -- CrossCat enables rapid screening
- 4-chain Bayesian Model Averaging on consumer GPU (GTX 1650, 4GB VRAM)
- 3-step reproducible pipeline: fetch (2 min) + preprocess (45s) + predict (94 min)

Data: Materials Project v2025.09.25 | Compute: NVIDIA GTX 1650 (4GB) | Total runtime: ~4 hours

1. The Problem: DFPT is Expensive

The Materials Project database contains 154,879 materials, but only 7,327 (4.7%) have dielectric constants computed via Density Functional Perturbation Theory (DFPT). DFPT requires 5-10x more compute than standard DFT: higher k-point densities (3,000/atom vs 1,000), tighter convergence, and 600 eV energy cutoffs. At current rates, computing dielectric constants for all remaining materials would require millions of additional CPU-hours.

We demonstrate that CrossCat -- a Bayesian nonparametric model -- can predict ionic dielectric constants at $R^2=0.81$ from cheap structural features alone, producing 49,566 high-confidence predictions at 99% CI without any DFPT calculation.




Figure 1: Dataset coverage. Of 154,879 Materials Project materials, only 7,327 have DFPT dielectric data (training set). CrossCat predicts dielectric for 147,552 new materials, with 49,566 meeting the 99% CI threshold.

2. Approach

CrossCat jointly discovers: (1) which properties are statistically dependent, and (2) which materials behave similarly. All parameters are integrated out via conjugate priors -- only structural assignments are sampled via Gibbs MCMC. This gives principled uncertainty estimates for free. JAX-CrossCat achieves 10-100x speedup via JIT compilation and vectorized operations.

- **Multi-chain inference:** 4 independent MCMC chains, $R_{\text{hat}}=1.007$ (converged)
- **Bayesian Model Averaging:** Predictions averaged across all 4 chains for stability
- **Calibrated uncertainty:** 99% CI from cross-chain std achieves 99.7% coverage on holdout

3. Structure Discovery: 5 Property Groups

CrossCat discovered 5 independent property groups -- consistent across all 4 MCMC chains. This is the result no supervised method can produce.

- **View 0 -- Structural/Thermodynamic (12 cols, 9 clusters)**

Band gap, formation energy, E above hull, stability, density, volume, nsites, nelements, crystal system, electronegativity, ionic radius, Laue class

- **View 1 -- Electronic/Mechanical (7 cols, 6 clusters)**

Metallicity, electronic dielectric, bulk/shear modulus, Poisson ratio, magnetization, magnetic ordering

- **View 2 -- Ionic Dielectric Pair (2 cols, 4 clusters)**

Ionic dielectric + total dielectric (tightly correlated)

- **View 3 -- Piezoelectric (1 col, 4 clusters)**

Piezo e_{ij_max} (independent singleton)

- **View 4 -- Elastic Anisotropy (1 col, 5 clusters)**

Universal anisotropy (independent singleton)

Key Physical Insight:

Ionic/total dielectric separated from electronic dielectric into different views. This reflects distinct physics: ionic dielectric depends on lattice dynamics (phonons), while electronic dielectric depends on band structure. CrossCat discovered this separation without any physics knowledge -- a validation of the model's structure discovery capability.




Figure 2: Z-matrix (dependence probability). Clear block structure: structural/thermodynamic (View 0), electronic/mechanical (View 1), and ionic dielectric pair (View 2).

4. Ground Truth Validation (99% CI)

Before predicting for new materials, we validate on the 7,327 training materials with known DFT dielectric constants. 10% of observed ionic dielectric values are masked and predicted using the remaining features.

Credible Interval Calibration:

CI Level	Expected	Actual	Status
90%	90%	95.6%	Conservative
95%	95%	98.4%	Conservative
99%	99%	99.7%	Near-perfect

The model is slightly conservative at all CI levels -- it rarely gives overconfident predictions. For the screening use case ('should I spend 5-10x compute on DFPT for this material?'), conservative uncertainty is preferable to optimistic uncertainty.




Figure 3: Holdout parity plot with 99% credible intervals. R²=0.81, 99% CI coverage=99.7%. Predictions cluster tightly around the diagonal across 3 orders of magnitude.

5. Predicting Dielectric for 49,566 Materials

Using the validated model, we predict ionic dielectric constants for 147,552 Materials Project materials that lack DFPT data. Each prediction uses Bayesian Model Averaging across 4 MCMC chains (n_samples=500 per chain). After filtering by 99% CI relative precision (CI width < prediction magnitude), 49,566 materials have high-confidence predictions.

Prediction Summary:

Total predictions	147,552 materials
High-confidence (99% CI)	49,566 materials (33.6%)
Ionic dielectric range	4.8 -- 80.6 (log1p scale)
Mean cross-chain std	0.27 (4 chains agree closely)
Mean confidence	0.795
Method	4-chain BMA, 500 samples/chain
Compute time	94 minutes (GTX 1650, 4GB VRAM)




Figure 4: Left: predicted ionic dielectric distribution matches DFT observed. Right: confidence distribution -- 99% CI subset (green) has high confidence, filtered materials (coral) have wide cross-chain disagreement.

6. Screening Candidates: Highest Predicted Dielectric

The 99% CI subset enables targeted DFPT calculations: instead of computing dielectric constants for all 147K materials, experimentalists can prioritize the top candidates identified by CrossCat. Each prediction includes uncertainty bounds to quantify confidence in the recommendation.




Figure 5: Top 30 materials by predicted ionic dielectric (99% CI). Blue circles = BMA prediction, red bars = 99% credible interval. Tight CI bars indicate high cross-chain agreement.




Figure 6: BMA quality -- prediction magnitude vs cross-chain uncertainty. Green = high confidence, red = low confidence. The 49,566 high-confidence materials (green cluster) have low cross-chain std.

7. Baseline Comparison

CrossCat achieves 88% of Random Forest R^2 on ionic dielectric (0.81 vs 0.92) while providing capabilities no supervised model can match:

- Structure discovery: 5 physically meaningful property groups (unique to CrossCat)
- Calibrated uncertainty: 99.7% coverage at 99% CI (RF provides no uncertainty)
- No per-target training: one model predicts all 23 properties simultaneously
- Native mixed types: continuous + binary + categorical + ordinal in one model
- Handles arbitrary missingness: 43.5% NaN in new materials handled natively
- Anomaly detection with attribution: identifies unusual property combinations




Figure 7: R^2 comparison on 10% holdout. RF wins on raw accuracy (green) but requires per-target training and provides no uncertainty. CrossCat (blue) is a general-purpose model that also discovers structure and quantifies uncertainty.

The comparison is inherently asymmetric: RF is a specialized point predictor; CrossCat is a general-purpose probabilistic model. CrossCat's value is not replacing RF -- it's providing structure discovery and calibrated uncertainty that RF cannot, while achieving competitive accuracy.

8. Summary of Contributions

- **49,566-Material Predicted Dielectric Dataset**

High-confidence ionic dielectric predictions for materials lacking DFPT data. 99% CI with 99.7% coverage. Publishable dataset for the materials community.

- **DFPT Screening Tool ($R^2=0.81$)**

Predicts from cheap structural features only -- no additional DFT needed. Saves 5-10x compute cost per material. 94-minute runtime on consumer GPU.

- **Structure Discovery (5 Views)**

Independently rediscovered phonon vs band structure physics separation. Consistent across all 4 MCMC chains. No supervision or domain knowledge used.

- **Calibrated Uncertainty (99.7% coverage)**

Conservative credible intervals at all levels (90/95/99%). Enables informed screening decisions: 'is this prediction trustworthy?'

- **Reproducible Pipeline**

3 scripts: fetch (2 min) + preprocess (45s) + predict (94 min). Runs on GTX 1650 (4GB). Code + checkpoint open-source on GitHub.

Future Work:

- **Publish predicted dataset:** Upload 49,566-material CSV to Materials Project community as a predicted dielectric dataset with uncertainty bounds.
- **Improve electronic dielectric:** Add band structure descriptors as features -- current R^2 limited by coarse proxies (is_metal, elastic moduli).
- **Publication:** Submit to npj Computational Materials or ICML/NeurIPS AI4Science. 5-view discovery is the hook, $R^2=0.81$ is the proof, 99.7% calibration is the practical value.

Resources

- GitHub: <https://github.com/sambhal-labs/jaxcross>
- Docs: <https://sambhal-labs.github.io/jaxcross/>
- Data: <https://materialsproject.org/> (API v2025.09.25)
- Notebook: [examples/materials_project_discovery_v2.ipynb](#)
- Pipeline: [examples/skills/{fetch,preprocess,impute}_*.py](#)